

- WHEN DOCKER ENDS  
CHEF BEGINS



Hello!

I AM GIOVANNI TORALDO

Open Source enthusiast with SuperCow Powers  
PHP/Java/whatever developer  
writer of the OpenNebula book  
Lead Developer @ ClouDesire

## ● WHAT IS CLOUDESIRE?

### ○ Application Marketplace

- Helps S/M software vendors
- For simple applications it can
  - provision VM
  - on multiple cloud providers
  - monitor resources
- For complex applications
  - expose REST API
- For everyone
  - manage subscriptions, billing, pay-per-use, invoicing, payments



1

WOULD YOU BE MY CONTAINER?

## ● DOCKER: WHAT IS IT?

- Enables software developers to
  - package an application
  - with all dependencies
  - runs it everywhere unchanged



## ● DOCKER: WHAT IS THE POINT?


- Enables system administrators to
  - simplify application deployment
  - ease scale-up & scale-down
  - processes separation





## ● DOCKER: UNDERLYING TECHNOLOGIES

- *Linux namespaces*
- *Control Groups (cgroups)*
- Layered filesystems
- *LXC (now libcontainer)*





## ● DOCKER: GLOSSARY

- **Image**: immutable snapshot of a container, push/pull repository
- **Container**: an instance launched from an image
- **Volume**: persistent writable area of a container
- **Registry**: repository of images (versioned via **tags**)
- **Dockerfile**: the descriptor from which an image is built

## ● DOCKER: HOW DO I RUN IT?

### ○ GNU/Linux

`wget -qO- https://get.docker.com/ | sh`

### ○ Windows

<https://github.com/boot2docker/windows-installer/releases/latest>

### ○ OSX

<https://kitematic.com/download/>

### ○ Hello world

*`$ docker run -ti ubuntu:14.04 /bin/bash`*

## ● DOCKER: WHAT HAPPENS UNDER THE HOOD?

- - **Pulls** the ubuntu image from registry
  - Creates a **new container**
    - Allocates a rw **filesystem**
    - Allocates a **network** interface (on a bridge)
    - Sets up network (**IP address**, dns..)
  - **Launch a process** in the container
  - Captures and provides **application output**

Container terminates when the process exit

## DOCKER: A SIMPLE DOCKERFILE

 Dockerfile

Raw

```
1 FROM ubuntu:14.04
2 MAINTAINER Giovanni Toraldo
3
4 ENV DEBIAN_FRONTEND noninteractive
5
6 RUN sed -i s/archive/it.archive/g /etc/apt/sources.list
7 RUN apt-get update
8
9 RUN apt-get install -y software-properties-common
10 RUN apt-add-repository ppa:chris-lea/node.js
11 RUN apt-get update
12 RUN apt-get install -y nodejs
13
14 RUN apt-get clean
15 RUN rm -rf /var/lib/apt/lists/*
16
17 RUN mkdir /var/www
18
19 ADD app.js /var/www/app.js
20
21 ENV PORT 8080
22 EXPOSE 8080
23
24 VOLUME ["/var/www/data"]
25
26 CMD ["/usr/bin/node", "/var/www/app.js"]
```

## ● DOCKER: STANDARD WORKFLOW

### ○ Build & push:

- `docker build -t gionn/nodejs-app:1.0.0 .`
  - a tagged image is generated
- `docker push gionn/nodejs-app:1.0.0`
  - publish to repository

### Pull & run:

- `docker pull gionn/nodejs-app:1.0.0`
  - fetch from repository
- `docker run gionn/nodejs-app:1.0.0`
  - run container from this image

Example gist: [link](#)

## ● DOCKER: ROUGH EDGE #1

○ Service in container A needs to talk to service in container B

Docker solution:

- Use Container Links

Reality:

- Works only on the same host
- Ordered sequence to boot-up
- Can't solve cyclic dependencies

## ● DOCKER: ROUGH EDGE #2

- My containerized application needs environment-dependant configurations

Docker solution:

- Inject environment variables

Reality:

- I need to fill YAML, XML, JSON complex structures

## ● DOCKER: ROUGH EDGE #3

○ I need to manage and upgrade a non-trivial number of containers on multiple hosts

Docker solution:

- Docker Swarm

Reality:

- currently in beta, not recommend for production



- DOCKER: RECAP


- So far so good?

Docker is a piece of cake for wrapping together the technologies of Linux containers, multi-layered filesystems and an image build system, in an unique tool easy and fast to use.

## ● DOCKER: RECAP

○ But what about the environment?

Being a (relatively) young project, the ecosystem of tools is pretty scattered and inconsistent.





2

WHO YOU GONNA CALL?

● WHO YOU GONNA CALL?

○ Probably someone has solved this kind of issues far time ago, even before Docker existed?

Those kind of problems are all about **configuration management** and **automation**.

So use the tools already available.

[Web](#)[News](#)[Images](#)[Shopping](#)[Maps](#)[More ▾](#)[Search tools](#)

About 2,620,000 results (0.50 seconds)

## SaltStack automation for CloudOps, ITOps & DevOps at scale

[saltstack.com/](https://saltstack.com/) ▾

SaltStack systems & **configuration management** software delivers fast & scalable event-driven **infrastructure automation** & predictive cloud orchestration.

## Configuration Management | Puppet Labs

<https://puppetlabs.com/solutions/configuration-management> ▾ Puppet ▾

Puppet Enterprise is IT **automation** software that makes it easy for systems administrators to provision, **configure**, and **manage infrastructure** throughout its ...

## What is Puppet? | Puppet Labs

[puppetlabs.com/puppet/what-is-puppet](https://puppetlabs.com/puppet/what-is-puppet) ▾ Puppet ▾

Puppet is a **configuration management** system that allows you to define the state of your IT **infrastructure**, then automatically enforces the correct state. ... You'll find more than 2,700 pre-built modules to **automate** common systems ...

## Configuration management - Wikipedia, the free encyclopedia

[en.wikipedia.org/wiki/Configuration\\_management](https://en.wikipedia.org/wiki/Configuration_management) ▾ Wikipedia ▾

**Configuration management** (CM) is a systems engineering process for establishing ... management methodology, COBIT, Information Technology **infrastructure** ...

## Chef | IT automation for speed and awesomeness | Chef

<https://www.chef.io/chef/> ▾

With Chef, you can **automate** how you build, deploy, and **manage your infrastructure**. ... Chef server stores your recipes as well as other **configuration data**.

## ● CHEF

○ Chef enables you to:

- **Version** your **infrastructure** on SCM, build an artifact
- Apply testing, CI, CD to infrastructure
- Keep it aligned with your software
- Automation via **repeatable** actions (e.g. click to deploy)

- CHEF: THE TOOLS

- Everything you need in a single package:

<https://downloads.chef.io/chef-dk/>

For (automated) testing

<https://www.vagrantup.com>

<https://www.virtualbox.org>

- CHEF: EVERYTHING IS IN A REPOSITORY

- The chef-repo is a standard repo layout and contains:

- Cookbooks
- Environments
- Data bags
- Roles



## ● CHEF: WHAT IS A COOKBOOK

○ Each cookbook is coupled with a service (e.g. mysql).

Contains:

- **Attributes:** they are like global variables (e.g. version to install)
- **Recipes:** an atomic unit of configuration
- **Templates:** patterns to generate real files, filled with data
- **Files:** static configuration

## CHEF RECIPES

Each recipe contains **behaviour** expressed by **resources** (and Ruby code)

```
user_name = 'gionn'
user user_name do
  supports :manage_home => true
  uid 1000
  gid 'users'
  home "/home/#{user_name}"
  shell '/bin/bash'
  password '$1$JJsvHslV$szsCjVEroftprNn4JHtDi'
end
```

## ● CHEF COMPONENTS

- The remaining components:
  - **Environments:** contains common attributes for a group of nodes
  - **Roles:** contains attributes for nodes sharing a particular behaviour
  - **Data bags:** general-purpose JSON data, optionally encrypted, usually to store credentials

## ● COOKBOOKS FOR EVERY NEEDS

○ All cookbooks are usually hosted on GitHub

- Maintained by Opscode

<https://github.com/opscode-cookbooks>

- by vendors

<https://github.com/elastic/cookbook-elasticsearch>

- by the community

<https://supermarket.chef.io>

Community Stats (07/04/2015)

2,120 Cookbooks ~ 62,086 Chefs

3

ONE DOES NOT SIMPLY COOK A WHALE

## ● DOCKER COOKBOOK

○ <https://github.com/bflad/chef-docker>

- Install docker daemon on supported platforms
  - Ubuntu/Debian
  - RHEL/CentOS/Fedora
  - Amazon Linux
- Expose attributes for fine-tuning (e.g. TLS certificates, DNS)
- Manage images & containers lifecycle via ad-hoc resources



This repository Search

Explore Gist Blog Help



glonn



bflad / chef-docker

Watch 37

Unstar 363

Fork 167

Contributors

Commits

Code frequency

Punch card


Network

Members

May 5, 2013 – Apr 4, 2015

Contributions to master, excluding merge commits


Contributions: Commits



bflad

374 commits / 10,019 ++ / 3,522 --

#1



tduffield

51 commits / 1,900 ++ / 758 --


#2



jperville

22 commits / 346 ++ / 72 --

#3



jayofdoom

13 commits / 78 ++ / 49 --

#4



## ● PROBLEM #0: IMAGE DISTRIBUTION / RUN CONTAINER

○

```
docker_image 'registry:0.9.0' do
  action :pull
  notifies :redeploy, 'docker_container[registry]', :immediately
end
```

```
docker_container 'registry' do
  container_name 'registry'
  image 'registry:0.9.0'
  detach true
  port '5000:5000'
  volume '/srv/registry:/tmp/registry'
  env 'SETTINGS_FLAVOR=local'
  env 'SEARCH_BACKEND=sqlalchemy'
  action :run
end
```



## ● PROBLEM #0: IMAGE DISTRIBUTION / RUN CONTAINER

○ Check that docker has such tagged image and if not:

- Download that version
- Stop existing container (if any)
- Run new container
- Raise error if anything goes wrong

## ● PROBLEM #1: MY APPLICATION NEEDS CONFIGURATION

○ Populate configuration files with proper values (and automatically restart on changes)

```
template node['tomcat']['host'] + '/cmw.xml' do
  source 'tomcat/tomcat-context.xml.erb'
  variables(
    :resource_name => 'jdbc/datasource',
    :username       => node['cloudesire']['name'],
    :password       => node['cloudesire']['pass'],
    :url            => node['cloudesire']['url'],
  )
  notifies :redeploy, 'docker_container[cmw]'
end
```

## PROBLEM #1: MY APPLICATION NEEDS CONFIGURATION

### tomcat-context.xml.erb

```
<Context>
  <Resource auth="Container"
    driverClassName="org.postgresql.Driver"
    initialSize="<%= @node['cloudesire']['backend']['db_pool']['min'] %>"
    maxActive="<%= @node['cloudesire']['backend']['db_pool']['max'] %>"
    maxIdle="<%= @node['cloudesire']['backend']['db_pool']['idle'] %>"
    maxWait="<%= @node['cloudesire']['backend']['db_pool']['wait'] %>"
    name="<%= @resource_name %>"
    url="<%= @url %>"
    username="<%= @username %>"
    password="<%= @password %>"
    type="javax.sql.DataSource"
    validationQuery="select 1"
    testOnBorrow="true"
  />
</Context>
```

## ● PROBLEM #1: MY APPLICATION NEEDS CONFIGURATION

- Templates consist of:
  - an .ERB template
  - a template resource declared in a recipe

The template is evaluated using the variables passed directly or via the global node object.

## ● PROBLEM #1: MY APPLICATION NEEDS CONFIGURATION

### ○ Inject a single file or entire folders

```
dst = node['tomcat']['base'] + '/conf/Catalina/' +  
      'localhost/cmw.xml'  
  
docker_container 'cmw' do  
  image image_name  
  container_name 'cmw'  
  detach true  
  
  env LOG debug  
  volume [  
    "#{node['tomcat']['host']}/cmw.xml:#{dst}",  
    "/etc/cloudesire:/etc/cloudesire"  
  ]  
end
```

- PROBLEM #1: MY APPLICATION NEEDS CONFIGURATION

- Docker permits defining **volumes** to be used for persistent data (e.g. database files), but may be used to inject configurations into the container at runtime.

Definitely **avoid** the needs of image rebuilding to adjust a setting.

## ● PROBLEM #2: CONTAINERS RUNNING ON MULTIPLE HOSTS

○ Each node has its own **run\_list**, defining which recipes should be executed (in JSON):

```
{
  "run_list": [
    "cd-infrastructure::docker-cmw",
    "cd-infrastructure::docker-deployer",
    "cd-infrastructure::docker-monitor",
    "cd-infrastructure::docker-logger"
  ],
  "cloudesire": {
    "key": "value"
  }
}
```

## ● PROBLEM #2: CONTAINERS RUNNING ON MULTIPLE HOSTS

○ Same recipe on different nodes  
(attributes may change)

### **node1.json**

```
{
  "run_list": [
    "cd-infrastructure::docker-cmw",
    "cd-infrastructure::docker-logger"
  ]
}
```

### **node2.json**

```
{
  "run_list": [
    "cd-infrastructure::docker-deployer",
    "cd-infrastructure::docker-monitor",
    "cd-infrastructure::docker-logger"
  ]
}
```



● MAY NOT BE GOLD BUT IT'S A START FOR SURE!

○ It's easy to getting started with Chef by using kitchen-ci or plain Vagrant:

- Initialize a chef repo
- Create a new cookbook
- Start hacking
- Play on kitchen-ci or vagrant
- Repeat last 2

## KITCHEN.YML EXAMPLE

\$ kitchen converge

```
driver:
  name: vagrant

provisioner:
  name: chef_solo

platforms:
  - name: ubuntu-1404

suites:
  - name: default
    run_list:
      - recipe[mycookbook::docker-whatever]
    attributes: { foo: "bar" }
```

- READY TO USE CHEF REPOSITORY

- A starting repository for aspiring whale cooks:

<https://github.com/gionn/cooking-docker>

## ● DOCKER APPENDIX: GOLDEN RULES

- - Only one process per Image
  - No embedded configuration
  - No, you don't need SSH
  - No, you don't need syslog
  - No, you won't touch a running container to adjust a thing
  - No, you will not use a community-contributed image without looking at what it do

**Thanks!**

**ANY QUESTIONS?**

You can find me at:

@gionn on twitter / github

me@gionn.net

giovanni.toraldo@cloudesire.com